

23CSE352: Big Data Analytics

Project Review 1 Report

Big Data Analytics of Crime Patterns using Hadoop MapReduce and Apache Hive

Team Members:

Sheela Akshar Sakhi

Nishanth S Gowda

August 13, 2026

Contents

1	Introduction	3
2	Problem Statement	3
3	Motivation	3
4	Dataset	3
5	Architecture	4
6	MapReduce Problem	4
7	Mapper logic explanation	4
8	MapReduce source code	4
9	Reducer Logic	5
10	Reducer code	5
11	Input and Output	5
12	Screenshots of execution	5
12.1	HDFS Storage	6
12.2	MapReduce Execution	6
13	Output screenshots	6
14	Hive analysis questions	7

15 Each query executions	7
16 Screenshot Query result	8
17 Conclusion	8

1 Introduction

This project aims to process and analyze crime data from the City of Chicago using Big Data technologies. With millions of records detailing various aspects of crime occurrences, this data requires distributed processing tools like Hadoop MapReduce and Apache Hive. By leveraging these frameworks, we extract meaningful insights, such as frequency of crime types, high-crime areas, and historical trends, transforming raw data into actionable intelligence.

2 Problem Statement

Law enforcement agencies and city planners require accurate and fast analysis of crime records to optimize resource allocation and enhance public safety. Processing large-scale crime datasets using traditional sequential systems is inefficient and time-consuming. The problem addressed in this project is to implement a scalable pipeline that efficiently counts crime frequencies and performs complex SQL-like analytics on a large real-world dataset.

3 Motivation

The motivation for this project stems from the critical need for data-driven decision-making in public safety. Analyzing crime patterns helps in understanding the socio-economic impact of crimes and assists police departments in predictive policing and strategic deployment. Utilizing Big Data technologies ensures that the analysis can scale with continuously growing datasets.

4 Dataset

Name: Chicago Crimes (2001 to Present)

Source: <https://data.cityofchicago.org/Public-Safety/Crimes-2001-to-Present/ijzp-q8t2>

Description: A comprehensive real-world dataset containing incident reports of crimes occurred in the City of Chicago. We have fetched a subset of 10,000 recent records via the official API to satisfy the project constraints. The key attributes are summarized below:

Attribute Category	Fields
Identifiers	id, case_number, date
Crime Typology	primary_type, description
Location Data	location_description, district, ward, community_area
Boolean Flags	arrest, domestic

Table 1: Dataset Attributes Overview

5 Architecture

The project follows a standard Big Data processing pipeline:

1. **Data Ingestion:** The dataset is downloaded as a CSV file, cleaned to remove embedded newlines, and uploaded to the Hadoop Distributed File System (HDFS).
2. **MapReduce Processing:** A Java-based MapReduce job is executed on Hadoop to parse the data in parallel and aggregate crime counts.
3. **Hive Analytics:** The dataset is loaded into Apache Hive with an `OpenCSVSerde`. 10 complex queries are executed to perform advanced data summarization.

6 MapReduce Problem

The MapReduce problem we are solving is **Category-wise Statistics**: determining the frequency of each primary crime type (e.g., THEFT, BATTERY) across the entire dataset.

7 Mapper logic explanation

The Mapper takes a line of the CSV file as input. It skips the header row. It then splits the line using a regular expression that correctly handles commas inside double quotes. The `primary_type` attribute is extracted from the 6th column, cleaned of quotes, and emitted as the intermediate key with a value of 1.

8 MapReduce source code

Below is the Java MapReduce code containing the Mapper (and Driver) classes:

```
1 package bigdata;
2
3 import java.io.IOException;
4 import org.apache.hadoop.conf.Configuration;
5 import org.apache.hadoop.fs.Path;
6 import org.apache.hadoop.io.IntWritable;
7 import org.apache.hadoop.io.Text;
8 import org.apache.hadoop.mapreduce.Job;
9 import org.apache.hadoop.mapreduce.Mapper;
10 import org.apache.hadoop.mapreduce.Reducer;
11 import org.apache.hadoop.mapreduce.lib.input.FileInputFormat;
12 import org.apache.hadoop.mapreduce.lib.output.FileOutputFormat;
13
14 public class CrimeTypeCount {
15
16     public static class CrimeMapper extends Mapper<Object, Text, Text,
17         IntWritable> {
18         private final static IntWritable one = new IntWritable(1);
19         private Text crimeType = new Text();
20
21         public void map(Object key, Text value, Context context) throws
22             IOException, InterruptedException {
```

```

21         String line = value.toString();
22         if (line.startsWith("\"id\"")) return; // Skip header
23
24         String[] fields = line.split(",(?=(?:[^\"]*\"[^\"]*\")"
25         *"[^\"]*$)");
26         if (fields.length > 5) {
27             String type = fields[5].replaceAll("^\"|\"$", "").trim
28             ();
29             if (!type.isEmpty()) {
30                 crimeType.set(type);
31                 context.write(crimeType, one);
32             }
33         }
34     }

```

9 Reducer Logic

The Reducer receives the intermediate key (`primary_type`) and an iterable list of values (mostly 1s) emitted by the Mappers. It iterates through these values, summing them up to calculate the total frequency of each crime type, and emits the final result as (`CrimeType`, `TotalCount`).

10 Reducer code

```

1     public static class CrimeReducer extends Reducer<Text, IntWritable,
2     Text, IntWritable> {
3         private IntWritable result = new IntWritable();
4
5         public void reduce(Text key, Iterable<IntWritable> values,
6         Context context) throws IOException, InterruptedException {
7             int sum = 0;
8             for (IntWritable val : values) {
9                 sum += val.get();
10            }
11            result.set(sum);
12            context.write(key, result);
13        }
14    }

```

11 Input and Output

Input: A cleaned CSV file residing in HDFS at `/user/$USER/bigdata-project/crimes/chicago_crime.csv`

Output: A text file generated by Hadoop in `/user/$USER/bigdata-project/crime_output/` containing lines of tab-separated crime types and their total counts (e.g., `THEFT 2150`).

12 Screenshots of execution

(Replace the placeholder images with your actual screenshots)

12.1 HDFS Storage

```
14282717 NARCOTICS MANUFACTURE / DELIVER - CANNABIS OVER 10 GRAMS true
14282623 NARCOTICS POSSESS - HEROIN (TAN / BROWN TAR) true
14282521 NARCOTICS POSSESS - HEROIN (WHITE) true
14282606 NARCOTICS POSSESS - HEROIN (WHITE) true
hadoop@aksharsakhi-QEMU-Virtual-Machine: ~/Big-Data-Analytics-of-Crime-Patterns$ hdfs dfs -ls /user/hadoop/bigdata_project/crimes/
2026-08-13 16:30:20,448 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-ja
va classes where applicable
Found 1 items
-rw-r--r-- 1 hadoop supergroup 2484524 2026-08-13 15:50 /user/hadoop/bigdata_project/crimes/chicago_crimes_clean.csv
hadoop@aksharsakhi-QEMU-Virtual-Machine: ~/Big-Data-Analytics-of-Crime-Patterns$
```

Figure 1: Uploading dataset to HDFS and verifying storage.

12.2 MapReduce Execution

```
1 warning
Creating JAR file...
added manifest
adding: bigdata/(in = 0) (out= 0)(stored 0%)
adding: bigdata/CrimeTypeCount.class(in = 1695) (out= 923)(deflated 45%)
adding: bigdata/CrimeTypeCount$CrimeReducer.class(in = 1765) (out= 746)(deflated 57%)
adding: bigdata/CrimeTypeCount$CrimeMapper.class(in = 2061) (out= 909)(deflated 55%)
Removing old HDFS output directory if exists...
2026-08-13 15:52:40,240 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-ja
va classes where applicable
Running MapReduce job...
2026-08-13 15:52:41,197 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-ja
va classes where applicable
2026-08-13 15:52:41,674 INFO client.DefaultNoHARMFalloverProxyProvider: Connecting to ResourceManager at localhost/127.0.0.1:8
032
2026-08-13 15:52:42,071 WARN mapreduce.JobResourceUploader: Hadoop command-line option parsing not performed. Implement the To
ol interface and execute your application with ToolRunner to remedy this.
2026-08-13 15:52:42,105 INFO mapreduce.JobResourceUploader: Disabling Erasure Coding for path: /tmp/hadoop-yarn/staging/hadoop
/.staging/job_1786616274630_0001
2026-08-13 15:52:42,357 INFO input.FileInputFormat: Total input files to process : 1
2026-08-13 15:52:42,826 INFO mapreduce.JobSubmitter: number of splits:1
2026-08-13 15:52:43,385 INFO mapreduce.JobSubmitter: Submitting tokens for job: job_1786616274630_0001
2026-08-13 15:52:43,385 INFO mapreduce.JobSubmitter: Executing with tokens: []
2026-08-13 15:52:43,564 INFO conf.Configuration: resource-types.xml not found
2026-08-13 15:52:43,564 INFO resource.ResourceUtils: Unable to find 'resource-types.xml'.
2026-08-13 15:52:44,018 INFO impl.YarnClientImpl: Submitted application application_1786616274630_0001
2026-08-13 15:52:44,050 INFO mapreduce.Job: The url to track the job: http://aksharsakhi-QEMU-Virtual-Machine:8088/proxy/appli
cation_1786616274630_0001/
2026-08-13 15:52:44,050 INFO mapreduce.Job: Running job: job_1786616274630_0001
2026-08-13 15:52:50,159 INFO mapreduce.Job: Job job_1786616274630_0001 running in uber mode : false
```

Figure 2: Executing the MapReduce Job on the Hadoop Cluster.

13 Output screenshots




Figure 3: MapReduce Output showing crime counts.

14 Hive analysis questions

The following 10 meaningful queries were executed in Apache Hive on the dataset to extract business intelligence. Table 2 maps the rubric requirements to the implemented queries.

Requirement	Query Implementation Focus
SELECT	Retrieving basic incident records
WHERE	Filtering specific crime types (Narcotics) and arrest statuses
ORDER BY	Sorting homicides chronologically
GROUP BY	Aggregating incidents by location description
HAVING	Filtering for dominant crime categories (count > 500)
COUNT	Counting occurrences per category or location
SUM	Calculating absolute total arrests per district
AVG	Determining average domestic crime rates per beat
MAX & MIN	Pinpointing districts with highest and lowest crime volumes
JOIN	Merging crimes with community area names for readability

Table 2: Mapping of Hive SQL Requirements to Queries

15 Each query executions

1. SELECT - Retrieve the first 10 crime records.

```

1 SELECT id, primary_type, description, crime_date
2 FROM crimes LIMIT 10;

```

2. WHERE - Find all incidents where an arrest was made for narcotics.

```

1 SELECT id, primary_type, description, arrest
2 FROM crimes WHERE primary_type = 'NARCOTICS' AND arrest = 'true' LIMIT
   10;

```

3. ORDER BY - List recent homicides ordered by date.

```

1 SELECT case_number, primary_type, crime_date
2 FROM crimes WHERE primary_type = 'HOMICIDE'
3 ORDER BY crime_date DESC LIMIT 10;

```

4. GROUP BY & COUNT - Find the number of crimes per location description.

```

1 SELECT location_description, COUNT(*) AS crime_count
2 FROM crimes GROUP BY location_description
3 ORDER BY crime_count DESC LIMIT 15;

```

5. HAVING - Find primary crime types that occurred more than 500 times.

```

1 SELECT primary_type, COUNT(*) as total
2 FROM crimes GROUP BY primary_type
3 HAVING COUNT(*) > 500 ORDER BY total DESC;

```

6. SUM - Calculate the total number of arrests made per district.

```

1 SELECT district, SUM(CASE WHEN arrest = 'true' THEN 1 ELSE 0 END) AS
   total_arrests
2 FROM crimes WHERE district IS NOT NULL AND district != ''
3 GROUP BY district ORDER BY total_arrests DESC LIMIT 10;

```

7. AVG - Calculate the average number of domestic crimes per beat.

```

1 SELECT beat, AVG(CASE WHEN domestic = 'true' THEN 1.0 ELSE 0.0 END) AS
   avg_domestic
2 FROM crimes WHERE beat IS NOT NULL AND beat != ''
3 GROUP BY beat ORDER BY avg_domestic DESC LIMIT 10;

```

8. MAX - Find the district with the maximum reported cases.

```

1 SELECT district, COUNT(*) AS district_total
2 FROM crimes WHERE district IS NOT NULL AND district != ''
3 GROUP BY district ORDER BY district_total DESC LIMIT 1;

```

9. MIN - Find the district with the minimum reported cases.

```

1 SELECT district, COUNT(*) AS district_total
2 FROM crimes WHERE district IS NOT NULL AND district != ''
3 GROUP BY district ORDER BY district_total ASC LIMIT 1;

```

10. JOIN - Join crimes and community_areas to display community names.

```

1 SELECT c.case_number, c.primary_type, a.area_name
2 FROM crimes c JOIN community_areas a ON (c.community_area = a.area_code
   ) LIMIT 15;

```

16 Screenshot Query result

17 Conclusion

This project successfully demonstrated the end-to-end processing of a real-world Big Data application. Using the Chicago Crimes dataset, we successfully ingested data into HDFS,


```
hadoop@aksharsakhi-QEMU-Virtual-Machine: ~/Big-Data-Analytics-of-Crime-Patterns
Initialization script completed
schemaTool completed
hadoop@aksharsakhi-QEMU-Virtual-Machine:~/Big-Data-Analytics-of-Crime-Patterns$ ./run_hive.sh
Starting Hive script execution...
This may take a few minutes.
Hive execution completed successfully.
Results saved to hive_output.txt.

=== Top 20 lines of Hive Output ===
14285893 BURGLARY BURGLARY FROM MOTOR VEHICLE 2026-08-04T00:00:00.000
14283703 CRIMINAL DAMAGE TO PROPERTY 2026-08-04T00:00:00.000
14283908 OTHER OFFENSE VIOLATE ORDER OF PROTECTION 2026-08-04T00:00:00.000
14283293 THEFT $500 AND UNDER 2026-08-04T00:00:00.000
14284321 CRIMINAL DAMAGE TO VEHICLE 2026-08-04T00:00:00.000
14283777 BATTERY AGGRAVATED DOMESTIC BATTERY - OTHER DANGEROUS WEAPON 2026-08-04T00:00:00.000
14284479 CRIMINAL DAMAGE TO VEHICLE 2026-08-04T00:00:00.000
14283403 DECEPTIVE PRACTICE STATE BENEFITS FRAUD 2026-08-04T00:00:00.000
14283429 MOTOR VEHICLE THEFT THEFT / RECOVERY - AUTOMOBILE 2026-08-04T00:00:00.000
14283519 CRIMINAL DAMAGE TO VEHICLE 2026-08-04T00:00:00.000
14283075 NARCOTICS POSSESS - COCAINE true
14283074 NARCOTICS POSSESS - CANNABIS 30 GRAMS OR LESS true
14283062 NARCOTICS MANUFACTURE / DELIVER - CANNABIS OVER 10 GRAMS true
14282837 NARCOTICS POSSESS - CRACK true
14283009 NARCOTICS POSSESS - CANNABIS MORE THAN 30 GRAMS true
14282739 NARCOTICS MANUFACTURE / DELIVER - CANNABIS OVER 10 GRAMS true
14282717 NARCOTICS MANUFACTURE / DELIVER - CANNABIS OVER 10 GRAMS true
14282623 NARCOTICS POSSESS - HEROIN (TAN / BROWN TAR) true
14282521 NARCOTICS POSSESS - HEROIN (WHITE) true
14282606 NARCOTICS POSSESS - HEROIN (WHITE) true
hadoop@aksharsakhi-QEMU-Virtual-Machine:~/Big-Data-Analytics-of-Crime-Patterns$
```

Figure 4: Screenshots showing successful execution of Hive queries.

executed a Java MapReduce program to calculate crime type frequencies, and leveraged Apache Hive to perform advanced SQL-based analytical queries. The pipeline provides a scalable foundation for public safety data analysis.